```
In [ ]: import pandas as pd
        import numpy as np
        import scipy.optimize as opt
        import statsmodels.api as sm
        from sklearn import preprocessing
        import matplotlib.pyplot as plt
        import seaborn as sns
        df=pd.read csv("/content/diabetes prediction dataset.csv")
        df.head()
Out[]:
           gender age hypertension heart_disease smoking_history
                                                                 bmi HbA1c_level blooc
        0 Female 80.0
                                 0
                                              1
                                                          never 25.19
                                                                              6.6
        1 Female 54.0
                                              0
                                                         No Info 27.32
                                 0
                                                                              6.6
        2
             Male 28.0
                                 0
                                                                              5.7
                                              0
                                                          never 27.32
        3 Female 36.0
                                 0
                                              0
                                                         current 23.45
                                                                              5.0
             Male 76.0
                                              1
                                                         current 20.14
                                                                              4.8
                                 1
        INFORMATION OF THE DATASET
In [ ]: df.info()
       <class 'pandas.core.frame.DataFrame'>
       RangeIndex: 100000 entries, 0 to 99999
       Data columns (total 9 columns):
        #
            Column
                                  Non-Null Count
                                                   Dtype
                                  100000 non-null object
        0
            gender
        1
            age
                                  100000 non-null float64
                                  100000 non-null int64
        2
            hypertension
            heart disease
                                  100000 non-null int64
        3
        4
            smoking history
                                  100000 non-null object
        5
            bmi
                                  100000 non-null float64
                                  100000 non-null float64
        6
            HbA1c level
        7
            blood_glucose_level 100000 non-null int64
                                  100000 non-null int64
            diabetes
       dtypes: float64(3), int64(4), object(2)
       memory usage: 6.9+ MB
```

In []: df.shape

Out[]: (100000, 9)

In []: df.describe()

Out[]:			age	hypertension	heart_disease	bmi	HbA1c_level	blood
	C	ount	100000.000000	100000.00000	100000.000000	100000.000000	100000.000000	
	m	nean	41.885856	0.07485	0.039420	27.320767	5.527507	
		std	22.516840	0.26315	0.194593	6.636783	1.070672	
		min	0.080000	0.00000	0.000000	10.010000	3.500000	
		25%	24.000000	0.00000	0.000000	23.630000	4.800000	
		50%	43.000000	0.00000	0.000000	27.320000	5.800000	
		75%	60.000000	0.00000	0.000000	29.580000	6.200000	
		max	80.000000	1.00000	1.000000	95.690000	9.000000	

In []: df.isnull()

Out[]:	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level
0	False	False	False	False	False	False	False
1	False	False	False	False	False	False	False
2	False	False	False	False	False	False	False
3	False	False	False	False	False	False	False
4	False	False	False	False	False	False	False
99995	False	False	False	False	False	False	False
99996	False	False	False	False	False	False	False
99997	False	False	False	False	False	False	False
99998	False	False	False	False	False	False	False
99999	False	False	False	False	False	False	False

100000 rows × 9 columns

Warning: total number of rows (100000) exceeds \max_{rows} (20000). Limiting to (20000) rows.

CHECKING OF DUPLICATE VALUES

In []: df.duplicated().sum()

Out[]: 3854

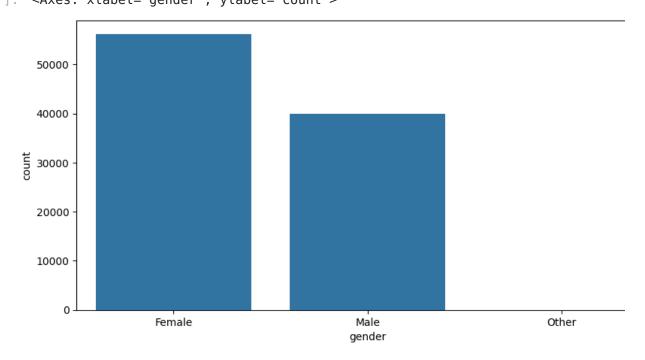
REMOVAL OF DUPLICATE VALUES

In []: df.drop_duplicates(inplace=True)


df.duplicated().sum()

Out[]: 0

GENDER DISTRIBUTION INFORMATION



DIFFERENT AGE GROUPS HAVING DIABETES



PEOPLE INVOLVED IN SMOKING

```
In [ ]: df['smoking_history'].value_counts()
```

```
Out[]:
                        count
        smoking_history
                 never 34398
                No Info 32887
                former
                         9299
                current
                        9197
             not current
                        6367
                        3998
                  ever
       dtype: int64
        PEOPLE SUFFERING FROM HEART DISEASE
In [ ]: df['heart_disease'].value_counts()
Out[]:
                     count
        heart_disease
                   0 92223
                   1
                      3923
       dtype: int64
        PEOPLE SUFFERING FROM HYPERTENSION
In [ ]: df['hypertension'].value_counts()
Out[]:
                     count
        hypertension
                  0 88685
                    7461
       dtype: int64
In [ ]: df.drop(columns=['hypertension'],inplace=True)
In [ ]: df.head()
```

```
Out[ ]:
                  age heart_disease smoking_history
                                                      bmi HbA1c_level blood_glucose_level
           gender
        0 Female 80.0
                                  1
                                               never 25.19
                                                                   6.6
                                                                                     140
        1 Female 54.0
                                  0
                                             No Info 27.32
                                                                                      80
                                                                   6.6
        2
             Male 28.0
                                               never 27.32
                                  0
                                                                   5.7
                                                                                     158
          Female 36.0
                                              current 23.45
                                                                   5.0
                                                                                     155
             Male 76.0
                                  1
                                             current 20.14
                                                                   4.8
                                                                                     155
In [ ]: x=df.drop(columns=['diabetes'])
        y=df['diabetes']
In [ ]: from sklearn.model_selection import train_test_split
        x_train, x_test, train_y, test_y = train_test_split(x, y, test_size=0.2, |
        train_y = np.nan_to_num(train_y)
In [ ]: df.head()
Out[]:
           gender age heart_disease smoking_history
                                                      bmi HbA1c_level blood_glucose_level
                                               never 25.19
        0 Female 80.0
                                  1
                                                                   6.6
                                                                                     140
        1 Female 54.0
                                  0
                                             No Info 27.32
                                                                   6.6
                                                                                      80
        2
             Male 28.0
                                  0
                                               never 27.32
                                                                   5.7
                                                                                     158
        3 Female 36.0
                                              current 23.45
                                                                   5.0
                                                                                     155
                                                                                     155
        4
             Male 76.0
                                              current 20.14
                                                                   4.8
                                  1
        LINEAR REGRESSION
In [ ]: from sklearn.model_selection import train_test_split
        from sklearn.linear model import LinearRegression
        # Instead of recreating dummy variables, select the relevant columns for
        x = df[['age', 'HbA1c level', 'blood glucose level']] # Features
        y = df['diabetes'] # Target
        x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2, |
        model = LinearRegression()
        model.fit(x_train, y_train)
        y pred = model.predict(x test)
In [ ]: model.score(x_train,y_train)
Out[]: 0.3294808573465473
In [ ]: model.score(x test,y test)
Out[]: 0.33316792290875474
        LOGISTIC REGRESSION
In [ ]: from sklearn.linear_model import LogisticRegression
        from sklearn.metrics import accuracy_score
        log reg = LogisticRegression()
        log_reg.fit(x_train, y_train)
```

```
Out[]:
        ▼LogisticRegression
       LogisticRegression()
In [ ]: log reg.score(x train,y train)
Out[]: 0.9571740600135212
In [ ]: log reg.score(x test,y test)
Out[]: 0.9551742069682787
        ACCURACY, PRECISION, RECALL, F1 SCORE, CONFUSION MATRIX
In [ ]: from sklearn.metrics import accuracy_score, precision_score, recall_score
        def evaluate model(model, X test, y test):
            y_pred = model.predict(X_test)
            print("Accuracy:", accuracy_score(y_test, y_pred))
            print("Precision:", precision_score(y_test, y_pred))
            print("Recall:", recall score(y test, y pred))
            print("F1 Score:", f1_score(y_test, y_pred))
            print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
        evaluate_model(log_reg, x_test, y_test)
       Accuracy: 0.9551742069682787
       Precision: 0.8731537793223284
       Recall: 0.5839628123184195
       F1 Score: 0.6998607242339833
       Confusion Matrix:
        [[17363 146]
        [ 716 1005]]
        DECISION TREE
In [ ]: from sklearn.tree import DecisionTreeClassifier # Import DecisionTreeClas
        dt = DecisionTreeClassifier()
        dt.fit(x train, y train)
Out[]:
        ▼ DecisionTreeClassifier
       DecisionTreeClassifier()
In [ ]: dt.score(x train, y train)
Out[]: 0.9720604295595194
In [ ]: dt.score(x test, y test)
```

```
Out[]: 0.9682267290691627
        ACCURACY, PRECISION, RECALL, F1 SCORE, CONFUSION MATRIX
In [ ]: from sklearn.metrics import accuracy_score, precision_score, recall_score
        def evaluate model(model, X test, y test):
            y_pred = model.predict(X_test)
            print("Accuracy:", accuracy_score(y_test, y_pred))
            print("Precision:", precision score(y test, y pred))
            print("Recall:", recall_score(y_test, y_pred))
            print("F1 Score:", f1_score(y_test, y_pred))
            print("\nConfusion Matrix:\n", confusion_matrix(y_test, y_pred))
        evaluate_model(dt, x_test, y_test)
       Accuracy: 0.9682267290691627
       Precision: 0.9497568881685575
       Recall: 0.6809994189424753
       F1 Score: 0.7932318104906937
       Confusion Matrix:
        [[17447
                   621
        [ 549 1172]]
        TESTING OF THE MODEL
In [ ]: import numpy as np
        from sklearn.preprocessing import StandardScaler # Import StandardScaler
        # Example new patient data (replace with actual values)
        new patient = np.array([[67,6.5,200]]) # Sample feature values
        scaler = StandardScaler()
        scaler.fit(x train)
        new patient = scaler.transform(new patient)
        # Predict the outcome
        prediction = model.predict(new patient)
        print("Diabetes Prediction:", "Diabetic" if prediction[0] == 1 else "Non-I
       Diabetes Prediction: Non-Diabetic
       /usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
       UserWarning: X does not have valid feature names, but StandardScaler was fitt
       with feature names
         warnings.warn(
       /usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
       UserWarning: X does not have valid feature names, but LinearRegression was fi
       with feature names
         warnings.warn(
        FOR MULTIPLE PATIENTS
```

```
In [ ]: import numpy as np
        from sklearn.preprocessing import StandardScaler # Import StandardScaler
        # Example new patient data (replace with actual values)
        new patient = np.array([[70,6.6,140],[73,6.5,200]]) # Sample feature value
        scaler = StandardScaler()
        scaler.fit(x train)
        # Scale input (if you used scaling)
        new patient = scaler.transform(new patient)
        # Predict the outcome
        prediction = model.predict(new patient)
        for i, pred in enumerate(prediction):
            print(f"Patient {i+1}: {'Diabetic' if pred == 1 else 'Non-Diabetic'}"
       Patient 1: Non-Diabetic
       Patient 2: Non-Diabetic
       /usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
       UserWarning: X does not have valid feature names, but StandardScaler was fitt
       with feature names
         warnings.warn(
       /usr/local/lib/python3.11/dist-packages/sklearn/utils/validation.py:2739:
       UserWarning: X does not have valid feature names, but LinearRegression was fi
       with feature names
         warnings.warn(
```